

Assignment 1

Business Continuity Management with focus on Disaster Recovery and Data Backup

Andreas Ntaflos* Andreas Ntaflos[†] Christoph Seidl[‡]

December 23, 2014

Organizational Aspects of IT Security
VU 2.0 188.312 2014W

In this document we examine the implementation of company-wide data backup and restore solutions from a Business Continuity Management approach, specifically in a Disaster Recovery context. We discuss relevant concepts such as Business Continuity Planning and Disaster Recovery Planning, key metrics such as Recovery Point Objective, Recovery Time Objective and others which define the requirements of backup and restore solutions and techniques, and how these metrics must be based on a Business Impact Analysis to be useful in the business context. We look at the risks associated with backup and restore from an information security point of view and how to control for these risks when implementing a backup solution.

After reviewing different, representative backup solutions and techniques and discussing their merits and drawbacks—again from an information security point of view—we highlight various challenges and potential problems a company faces when implementing backup solutions. Finally we discuss future developments and how backup and disaster recovery requirements may change over time.

*Matrikelnummer e0326302, daff@pseudoterminal.org

[†]Matrikelnummer e1328682, e1328682@student.tuwien.ac.at

[‡]Matrikelnummer e0427434, e0427434@student.tuwien.ac.at

1 Business Continuity Management

Business Continuity Management (BCM) is a strategic management process covering *Business Continuity Planning* (BCP) and *Disaster Recovery Planning* (DRP) to ensure critical business functions can continue during and after major incidents (flood, fire, earthquake, terrorism, vandalism, ...) with minimal loss of operations, systems and life.

BCM provides a framework for integrating resilience with the capability for effective responses that protects the interests of an organisation's key stakeholders. The main objective of BCM is to allow the organisation to continue to perform business operations under various conditions [4].

In BCM the two most important artifacts are the *Business Continuity Plan* (BCP) and the *Disaster Recovery Plan* (DRP). The BCP ensures that critical business functions can continue to operate during a disaster and its aftermath. This may include moving critical equipment to standby locations, bringing up emergency power generators, performing normally automated business operations manually or temporarily hiring additional personnel.

The DRP is usually a part of the BCP and is focused on the IT infrastructure itself. It defines how to prepare the IT infrastructure for a disaster, what to do in the event of a disaster and how to restore IT operations so that the business can go back to operating normally. The DRP's scope is thus much narrower and technology-focused than the BCP as a whole.

Disaster recovery can be quantified by four key metrics:

Maximum Tolerable Downtime (MTD) The maximum outage time of a critical business function or system that can be endured before the impact on the company becomes unacceptably severe

Recovery Point Objective (RPO) The amount of data loss, measured in time, that can be sustained from a disaster without causing irreparable damage to the company and its business functions

Recovery Time Objective (RTO) The time it may take to restore business functions after a disaster

Work Recovery Time (WRT) The time it may take to recover data and bring infrastructure back so that normal operation can continue

The MTD consists of the RTO and the WRT while the RPO exists independently of the MTD. Figure 1 shows the metrics in relation to each other. The MTD is also known as the *Maximum Tolerable Period of Disruption* (MTPD).

It should be noted that IT vendors of backup and DR solutions tend to conflate RTO and WRT into just RTO, which makes sense considering that a business function or process can hardly be declared "restored" when it is still missing its key business data. We will refer to RTO in the same manner in the remainder of this document.

The values for MTD, RPO and RTO usually differ between systems, business functions and processes. It is thus important not to define them ad-hoc and on a whim but on a solid understanding of the business functions in question. This understanding comes from the Business Impact Analysis.

2 Business Impact Analysis

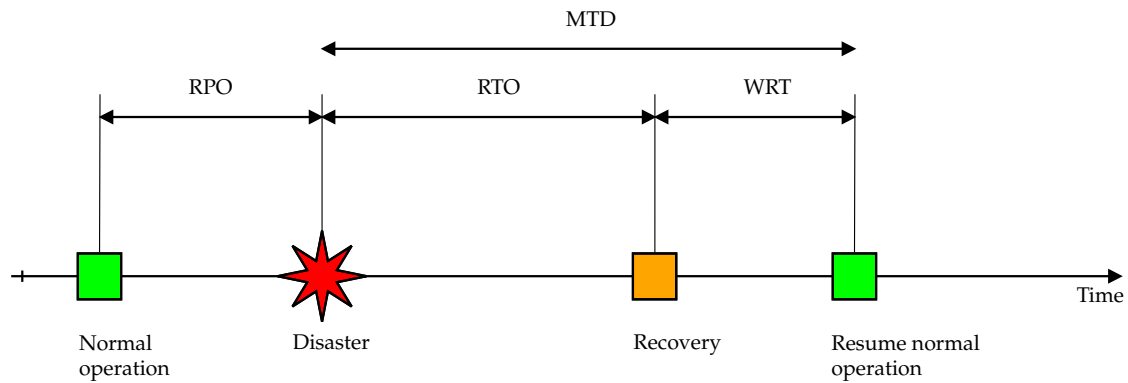


Figure 1: The disaster recovery metrics in relation to each other on a time line

2 Business Impact Analysis

The *Business Impact Analysis* (BIA) is a functional analysis that identifies the various business processes, functions, activities, resources and systems that define the business. It assigns them criticality levels and identifies threats and vulnerabilities for these functions and calculates the risk for each. The Business Impact Analysis is a key step in creating a disaster recovery plan and defining backup requirements.

In any reasonably sized organisation a BIA cannot be conducted by a single person or team without the help and input of key employees such as department heads and process owners. Thus the data for the BIA is gathered by means of interviews, questionnaires, workshops, existing documentation, etc. The result of the data gathering phase is a comprehensive list of business processes, functions and resources relevant to the organisation.

After all business processes and functions have been identified these resources need to be categorised by their criticality to the organisation. A resource is regarded as critical when its availability is imperative for the survival of the organisation and downtime of that resource is unacceptable. The most critical resources will receive the highest priority when recovering from a disaster while less- or non-critical systems will receive attention only after all critical systems have been dealt with. It is at this step that concrete values for MTD (MTPD), and later RPO and RTO of a critical resource can first be estimated.

MTD values can be categorised into *Nonessential*, *Normal*, *Important*, *Urgent* and *Critical*. The identified business resources are each assigned one of these categories and treated accordingly. The shorter the MTD, the more critical the resource, and the higher the priority it will be afforded when recovering.

When conducting the BIA it is important to consider interdependencies between business functions, resources or systems. Most business functions do not stand alone but depend on other business functions and systems, so when estimating the MTD these dependencies need to be taken into account.

The next step is threat and risk assessment. This includes identifying vulnerabilities,

3 Backup Requirements and Disaster Recovery Strategies

threats and to the most critical business resources, determining the impact each threat may have and the likelihood of its occurrence. Vulnerabilities in resources are single points of failure or security issues, threats may include sabotage, vandalism or theft but also natural disasters or major utilities outages. The risk to a business function or resource is thus calculated as $\text{Risk} = \text{Threat} \times \text{Impact} \times \text{Probability}$.

At this point it is important to note that while identifying and planning for specific threats with high impact or likelihood is prudent it is not useful to try to account and prepare for *all* possible threats. Instead the organisation should prepare for the loss of any or all business resources, regardless of specific threats. This will provide the organisation with the proper flexibility when the time of disaster recovery comes.

After determining the MTD for critical business functions taking into account their dependencies on other business functions and resources reasonable values for RPO and RTO can be defined. These are more meaningful than just the MTD in a disaster recovery context and are an important basis for each critical business function's disaster recovery strategy.

3 Backup Requirements and Disaster Recovery Strategies

Disaster recovery strategies for data and backup strategies are directly related. The backup and recovery strategy to choose for a business function or resource depends entirely on its defined RPO and RTO, determined during the Business Impact Analysis.

The backup frequency of a resource or system is defined by its RPO. When the RPO is two hours then a backup has to be done at least every two hours. When the RPO is a week then weekly backups are adequate. And when the RPO is zero then regular backups are not enough any more, because no potential data loss is allowed. In this case advanced backup/recovery techniques like synchronous mirroring off-site have to be employed.

What kind of backup technology and recovery strategy to use for a business resource is also directly influenced by the RTO. When the RTO of a resource is defined as four hours, but restoring data from tapes takes twelve then tape backup is not the right choice. In case of an RTO of zero the resource and its data must be available and functional throughout the disaster with no downtime at all. This calls again for advanced techniques like synchronous mirroring of data to a failover site that can take over operations from the primary site immediately. We will look at the various backup techniques and recovery strategies in this section.

According to [2] recovery strategies can be divided into seven tiers. The lower the tier the higher the values for RTO and RPO can be and the "cheaper" the recovery strategy can be. Figure 2 summarises the tiers and their properties. Other definitions for these seven (or more) tiers exist in literature; no standard definitions has been established so far.

4 Potential risks of backup solutions

Tier	Description
0	Do nothing, no off-site data
1	Data backup shipped off-site ("PTAM")
2	Data backup shipped off-site with hot-site
3	Electronic vaulting
4	Electronic vaulting to hot-site (active secondary site)
5	Two-site two-phase commit
6	Zero data loss

Figure 2: Seven-tier disaster recovery strategies with according RTOs and RPOs
(Source: [6])

4 Potential risks of backup solutions

The loss of the information system can be devastating for any company. As previously mentioned it will take some time to restart and restore the services. Incomplete data will slow down that process or even prevent it. One way to protect against threats is to use backup systems. Is it sufficient enough to have a backup system and just backup the data. Many companies have learned it the hard way, that performing backups on data is not enough. After incidents occurred they were not able to recover as they planned.

Backup solutions are not the universal remedy against any unforeseen incidents. They require intensive planning to prevent potential risks.

4.1 Availability

As previously mentioned one of the largest misconceptions about backup systems is that investing a large amount of money in the best solutions guarantee the possible recovery of the IT infrastructure. There are many factors to consider when building a backup solution for a company. One factor is the availability.

After any type of incidents it is not ensured that the backup system or even the backup data is available. There are two components that have to be considered when talking about availability of backup solutions. The hardware and the software.

The hardware is the component where the actual data is stored. The technologies in which the data is stored changed over the history of computers from punched cards to modern hard disks, but the issues remain the same. Hardware is susceptible to physical force. Hard disks are not very robust and even the smallest amount of force can damage it beyond repair. When choosing a backup solution this issue needs to be considered. While many of the incidents the company will face are rather small like power outages it is not impossible that the company will face larger incidents like natural disasters. When hit by an earthquake the chances are very high that if any type of damage occur on the information system the same damages will occur with the backup system if the company has not implemented any type of protection mechanism against it. Companies

4 Potential risks of backup solutions

are required to put their backup systems in different places than the information system. This should help to restore the system in case of natural disasters like earthquakes and flooding. The main disadvantage of this procedure is that putting the backup files somewhere else, that it can take some time to get the data or even risks that the data is damaged. Both advantages and disadvantages must be considered when making a decision on the type of the backup system.

Another problem with hardware components is that they are not reliable. While it is not very common nowadays, but the reliability of disks was a huge problem in the beginning of computer science. Under no circumstances hardware should be trusted. Even the backup systems need to have some type of mechanism to protect against hardware failures.

The software creates, manages, controls and restores backup files. This component is very underrated. It is great to store the backup data in different places than the information system, but it really does not help if there is no software to access the data that is needed to restore the services. The software should be available on different devices and not just run on the backup server. While this does not seem like a large problem enterprise backup solutions are complex systems. It is possible that the company that build the backup solution doesn't have any compatible versions of the program any more.

4.2 Integrity

The next big issue with backup systems is integrity. Even when every Bit was stored by the backup system, this does not guarantee that it is possible to restore the system. The key term for guaranteeing the recovery of the system is data integrity [8].

This term refers to maintaining and assuring the accuracy and consistency of data over its entire life cycle. When data is backed up, the system must store the data correctly. It can cause massive problems, if even a Bit is changed. For example it is a huge difference, if suddenly a comma is changed into a point. In America 123.000 is a complete different number than 123,000.

It is also important that the backup system does not only stores the data but stores the system state. There are some very important processes running on the information system. The backup system should store their state. For example the system should backup the status changes of shipping orders.

4.3 Confidentiality

Security is the final important aspect for a backup system. While on the first glance, this does not really look important for restoring the information system, but it is still very important for the management of the company. Security vulnerabilities can cause massive financial damages to the company.

Encryption The backup files must be encrypted. The backup files can be stored on different locations and are even transferred over the internet. It would be extremely easy,

5 Required Features

to simply start a man in the middle attack to get access to the data, or even gain access to the physical storage. The backup files often contain important data like customer informations and partner companies bank account numbers. It would be disastrous for the company if attackers would gain access to those informations over the company.

Access Control A important feature for backup solutions are access controls. Not only should the backup system encrypt the files, but the software that manages the files must be protected. The software should support state of the art access controls.

Version Management As a last line of defense the system should support a version management system. The system should store the version number, a version date and a hash value to protect the files against outside manipulations.

5 Required Features

The market for backup solutions is highly profitable. There are many companies, that offer different types of solutions for storing backup data. As discussed before, since there are many different types of risks, a certain amount of planning is required. So before choosing a backup solution certain topics need to be discussed:

- Definition of the data that needs to be backed up
- Definition of the frequency of backups
- Definition of the capacity that is needed to store the data
- Definition of legal compliance requirements
- Definition of the location
- Definition of proper security
- Definition of the proper tasks and jobs that are required
- Definition of proper system states
- Definition of a backup window

5.1 Data Recovery Plan

It is important for restoring the information system of ACME to backup following data:

- Warehouse information
- Production informations
- Business partner informations
- Customer informations
- Purchase order informations
- Shipment informations
- User accounts
- User email accounts
- Web services

5 Required Features

- User calendars
- System logs

5.2 Data Backup Plan

It is required, to backup the data that is defined in the data recovery plan after each working day. The purchase order informations should be stored every hour. The data should be stored in chronological order so that it is possible to access the backup data of certain days. The backups need to be stored for 7 years.

5.3 Location

A physical copy of the backup data should not be stored in the same building and the same district as the information system. While it is not recommended, to store the backup data in the same building the risk of not controlling the physical access and the availability after disaster counterbalance the risks of storing the data in the same building. To guaranty some sort of protection every Friday after closing time an encrypted backup of each information system should be transferred to either a partner company or if this is not possible an image should be send to offices in different cities or countries.

5.4 Backup Security

5.4.1 Integrity

It should not be possible to change the backup files. Backup files must have a unique version number and a creation date. To guaranty the safety of the files, every time a information system is performing a backup a hash value has to be created for the backup files.

5.4.2 Encryption Requirements

Every backup file should be encrypted. The encryption standards must be state of the art.

5.4.3 Access Control

Backup system must have an access control to prevent unauthorized access to the system. This access control must be separate from the access control of the regular information system.

5.5 Capacity

A recent survey has shown that ACME has a daily data usage of 1 Gigabyte. This includes information exchange between customers and sales agents, component supplier and

6 Types of backup solutions

the company and data created by employees. The total amount of data stored in the information system is 600 Gigabyte. Any backup solution should support at least a capacity of 40 Terabyte. Additionally it should support some types of compression algorithm.

5.6 Backup Window

Since ACME is a world wide operating company backups will occur at 2 am local time. During this time services will not be available. The backup process should only take a half hour. The systems should exit the backup process in the same state as they have entered it. Every process that was started before the backup should continue after it. Every process that was started during the backup should be stored and inserted after the system is started again.

6 Types of backup solutions

6.1 Traditional backup solutions

The most common type of enterprise backup solutions. Often it is a software suite and NAS servers. The software manages and controls the backup process of the data. The NAS servers, that should support at least Raid 10, store the data. The main advantage of this system set-up is the high level of security. The company can control the access of the data and can also control the encryption of the files. This level of security is also the main disadvantage of traditional backup solutions.

Customers bind themselves to certain companies and technologies. It is quite problematic to transfer the backup data from one backup solution to another. It also requires additional security mechanism to transfer the data to different locations. If the data is not transferred the chances are very high, that in case of natural disasters the backup data is lost or damaged. The transfer over the internet requires knowledge about security features and threats. The manual transfer requires a decent amount of times and also protection. Another problem with traditional backup solution is the scalability. The scaling to larger systems can lead to massive problems and can be quite expansive.

6.1.1 Products

IBM Tivoli Storage Manager A hard- and software solution by IBM. Price: About 120 US Dollar for each processor value unit.

6.2 Deduplication Backup solution

Unlike traditional backup solutions, deduplication solutions try to eliminate duplicate copies of repeating data. In the deduplication process unique elements of the data stream are stored. If the same pattern repeats itself in the data stream, the repeating pattern is deleted and a reference to the pattern is saved[7]. This mechanism decreases the amount

6 Types of backup solutions

of data that is stored and helps to save storage capacity. The disadvantage of this process is the speed. The deduplication process can be quite time extensive and the integrity of the data is not one hundred percent guaranteed. The storage process works in the same way as traditional backup solutions. After the data is deduplicated it is directly stored on different hardware components.

6.2.1 Products

EMC Data Domain Is considered by Gartner [5] as the best deduplication solution on the market. Price: 75000 US Dollar for the DD4200

DELL DR 6000 Dell offers their own hardware deduplication backup solution with the DR series. Price: 60000 US Dollar for the DR 6000

6.3 Cloud-based backup solutions

This phenomena started in the late '90s during the height of the dotcom bubble. The increase of the internet bandwidth made it possible to store large amount of data on external servers. While the storage of data over the internet is not new, it became a hype in recent years under the name cloud.

The cloud is a large number of remote server, that allow centralized storage of data. A cloud provider offers a large group of servers and a client software to store the data on the cloud. The backup data that is stored in the cloud, is constantly validated to guarantee whenever it is needed, the system is recoverable. One of the main advantages of cloud based backup solutions is the scalability. The service provider can easily expand the capacity for the customer. Another advantage is that the customer does not have to pay for hardware, but he has to pay monthly fees and for the amount of data he uploads.

While cloud based systems have their advantages they have serious disadvantages. The main disadvantages that should dis-encourage using cloud services is the security. It is not possible to control the access of the backup data on the cloud. The complete access control is managed by the cloud provider. Another huge disadvantage is the fact, that customer binds their company to the cloud provider. If the cloud service goes down the customer can not access the backup data. It is also a massive problem when the customer changes the cloud provider.

6.3.1 Products

Amazon S3 The market leader in cloud based storage solutions. It provides storage through Web services. Price: 0.15 US Dollar per Gigabyte

Microsoft Azure A Microsoft based cloud service and infrastructure. Price: 0.02 US Dollar per Gigabyte.

6.4 Hybrid Backup solution

Another approach for enterprise backup solutions are hybrid system. Hybrid system combine the advantages of different backup solutions and offer a complete package. This term was often used for physical backup solutions that offered deduplication functions. The advantage of this system was the decrease of the amount of the data. This increased the recovery time and also limited the needed capacity.

The true success story of hybrid backup solution started with cloud based backup solutions. While cloud based backup solutions have a decent amount of advantages, but the risks and disadvantages made companies cautious about using them. The solution for this problem was cloud based hybrid systems. The key words for this approach, are private and public cloud.

The private cloud grants access to the local information infrastructure with the same benefits and features of a public cloud. The private cloud can be managed by internal personal and is protected by the company security mechanism. The access of the cloud can be controlled and the data is stored locally.

The public cloud is what most people understand under cloud computing. As previously mentioned the public cloud is managed by the storage provider and the customer pays for storing and accessing the data in the cloud.

In hybrid cloud systems the customer can select, in which cloud his data is stored. With the private cloud the customer has the same advantages as with physical backup solutions. Additionally the customer can choose to save his data in the public cloud and have the same advantages as cloud based backup solutions.

6.4.1 Products

CommVault Simpana A hybrid cloud based storage system. The core component is the backup software that allows data to be encrypted in transit and at rest. It provides a enterprise wide view of the data and it also provides different audit systems. CommVault uses cloud services of different providers. Price: Up to 2000 US Dollar per agent (one server) and up to 2000 US Dollar per software module.

Datto Secures the backup data locally and sends the encrypted data to the public cloud. Additionally Datto devices work as restore hubs that restore everything from files to applications. Price: 3600 US Dollar plus 600 US Dollar for each 1 TB device.

7 Challenges and potential problems

7.1 Scalability

Scalability is the ability of a system to provide throughput in proportion to the data, and is limited only by the available hardware resources. A scalable system is one that can handle increasing numbers of requests without having its response time and throughput

7 Challenges and potential problems

negatively affected. The growth of computational power (usually by adding faster CPU, memory, etc.) within one operating environment is called vertical scaling. Horizontal scaling is adding and taking advantage of multiple systems that work together on a common problem in parallel.

The systems that will be running the backup and restore solutions should be horizontally and vertically scalable to account both for the increase of data in volume and plan for the future growth of ACME. Today, with the availability of large multi-core and large memory systems, there are more cases where a single machine can cover the scalability and performance goals. Still, there are several other factors to consider when choosing between the two options:

Continuous Availability/Redundancy One should always assume that failure is inevitable, and therefore having one big system is going to lead to a single point of failure. In addition, the recovery process is going to be fairly long which could lead to an extended down-time.

Cost/Performance Flexibility As hardware costs and capacity tend to vary quickly over time, flexibility is a very important factor in choosing the optimal configuration setup at any given time or opportunity to optimize cost/performance. If the backup and recovery system is designed for scale-up only, then you are pretty much locked into a certain minimum price driven by the hardware that you are using. In a competitive situation, the lack of flexibility could prove catastrophic for the business.

Continuous Upgrades Setting up the backup system as one one big unit is going to make it harder or even impossible to make future changes individually, without bringing the entire system down. In these cases it is perhaps better to decouple it into concrete sets of services that can be maintained independently.

Geographical Distribution There are cases where a backup system needs to be spread across data centers or geographical location to handle disaster recovery scenarios or to reduce geographical latency. In these cases you are forced to distribute it and the option of putting it in a single box doesn't exist.

With the availability of large multi-core machines at significantly lower costs, the question of vertical vs. horizontal scaling is more common than in earlier years, but it should not be considered as two distinct approaches that contradict one another, but instead as two complementing paradigms.

7.2 Interoperability

Interoperability is the ability of the backup solution to be able to support any system type operated by the business.

The biggest problem with a cross-platform backup solution is file system format incompatibilities. Files to be backed up have to be packaged in a format that won't cause their metadata to get altered or destroyed, so they will be restored as normal after a disaster scenario.

Further cases of interoperability issues can include the following:

7 Challenges and potential problems

- What happens if a file is locked or “busy” during backup?
- What happens if a number of files changes between the time a backup starts and ends?
- How are files larger than 2GB handled?

The types of systems that need to be supported will come up in the business evaluation plan, and it has to take into account the potential future growth and expansion. Just because a cloud solution is enough today, it doesn't mean that there won't be a need to purchase physical servers at some point in the future. Either way, we have to make sure the backup provider will scale with the company and enable our market growth, rather than hinder it.

The second aspect of interoperability is the ability of different version releases of the backup solution to communicate and work together in a distributed environment.

7.3 Vendor Lock-in

In the computer industry, both hardware and software, vendor lock-in can be used to describe situations in which there is a lack of compatibility or interoperability between equivalent components. Vendor lock-in is the restricted or proprietary use of a technology, solution or service developed by a vendor or vendor partner. This technique can be disabling and demoralizing because customers are effectively prevented from switching to alternate vendors.

Exit strategies, contract lock-in and data ownership (in cloud solutions) are the most important factors to consider when choosing a solution.

Especially for Cloud solutions, fear of vendor lock-in is clearly a legitimate concern. There are no widely adopted standards in place today to ensure that data can be freely moved among different cloud storage service providers' sites. It seems to be more difficult and more costly to transfer data out of a cloud storage repository than it is to upload it in the first place. The more data there is, the harder it is to move; and if the user wants to move the data to a new cloud provider the bandwidth costs could double. This inhibits opportunistic migration and puts customers in a weak negotiating position relative to their cloud storage service provider.

To reduce the likelihood of lock-in, and the cost and inconvenience that go along with it, the following guidelines are to be taken into consideration:

- Read the fine print of each provider's policies, and if necessary, ask them directly how they facilitate moving customer data out of their cloud storage repository. Given the amount of data and its expected growth over time, could the data be moved via the Internet back into your data center in a reasonable timeframe or to a different provider's site? As an alternative, if the volume of data is too large for digital transfer, can it be moved via a portable storage device? What's the process, timeframe and cost required for each of these approaches?
- Ask the provider whether they offer data migration tools or services to facilitate the movement of large amounts of data. For example, while most providers require that customers moving data between clouds first download data to an

7 Challenges and potential problems




Figure 3: Percentage of companies using multiple vendors

intermediate location (e.g., the customer's data center) and then re-upload the data to a new cloud. Some cloud gateway vendors are also making migration easier by integrating with different cloud storage APIs and then providing a standard file system interface to facilitate data migration between those clouds.

- Consider providers that have pledged to support emerging industry standards, such as the Cloud Data Management Interface (CDMI) standard created by the Storage Networking Industry Association (SNIA). The CDMI provides a standard functional interface that applications will use to create, retrieve, update and delete data elements in the cloud and, once adopted, will make it much easier to move data from one cloud to another.

7.4 One Size Fits All

That is a description for a product that would fit in all instances. The term has been extended to mean one style or procedure would fit in all related applications. It is not necessarily a negative term and always depends on the context of its use.

According to EMC's report on Global Data Protection Index , there exists a direct relationship between the number of vendors that a company relies on and figures such as IT spending, disruptions, data loss and the cost of downtime. The more vendors being used, the higher these damaging numbers climbed. Many companies may feel as if they have no choice but to use multiple vendors. Often times, backup support is needed for different operating systems. So, when a company needs support for both Windows and Linux, they'll often times look to the service of two vendors; one that supports Windows and one that supports Linux. The matter of fact is that backup solutions exist that provide support to more than one operating system. As shown in EMC's report, a single backup vendor is a financially preferable option for businesses needing protection over more than one operating system.

7 Challenges and potential problems




Figure 4:

7.4.1 Impact on Disruptions

Data loss and downtime are a company's worst nightmare. According to EMC, those organizations using one data protection vendor are the least likely to experience disruption.

There isn't a significant difference in disruption between companies using two vendors and those using three or more vendors. However, there is a significant leap in disruption when compared those using a single vendor. 54% of companies using three or more data protection solutions reported unplanned downtime, while only 42% of companies using a single vendor reported downtime. Similarly, 38% of those using three or more vendors experienced data loss compared to 24% of those who use single vendors.

7.4.2 Impact on Costs

Companies who use three or more data protection vendors are being hit with \$1.66 million in downtime costs. That's more than four times the amount that it costs those companies who are using just one data protection vendor (\$0.34 million).

The bottom line is that there is a much greater chance of experiencing disruption when using multiple data protection vendors, and that using multiple best-in-breed (highly specialized/cutting edge) solutions can cause more issues than they claim to solve.

8 Future Developments and Requirements

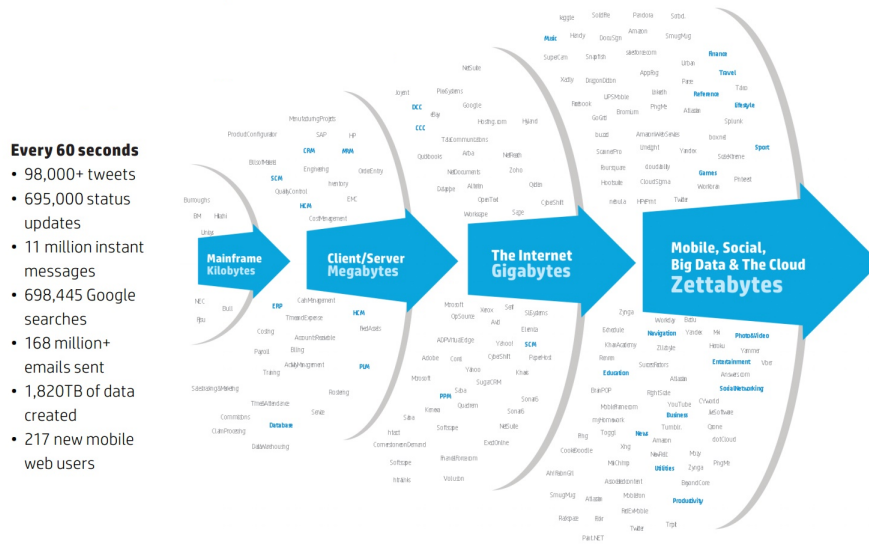


Figure 5: Data volume pace is accelerating

8 Future Developments and Requirements

8.1 Shift from traditional data structures to “Big Data”

As new information is created and handled within the enterprise, IT and its backup administrators and operators are faced with the challenge of protecting critical information assets at a much larger scale. Today’s information growth rates have IT organizations at their tipping points. They must “rethink” data protection and find the balance between serving the organization’s desire for big data, seeking more value from the information they create (through data mining and analysis), and the age old and essential requirement of protecting information from a disaster, corruption, or a logical or physical system failure.

Meeting the demands of big data backup require both retooling your approach to data capture and the use of optimization technologies. The topology of backup targets must evolve from the typical data-centric model to a more globally dispersed, flexible deployment approach to achieve better scale and data protection coverage.

Petabyte-size data stores can play havoc with backup windows, and traditional backup is not designed to handle millions of small files. Fortunately not all big data information needs to be backed up in the traditional way. Before considering how to protect your data, you should decide at what data needs to be protected. Machine-generated data – report data from a database, for instance – can be reproduced easier than it can be backed up and recovered. Compare the cost of protecting data with the cost of regenerating the data. In many instances, source data will need to be protected, but any post-acquisition processes may be cheaper to reproduce than the cost of protecting the

data after manipulation.

Although enterprises have become accustomed to rapid backup data growth, few are prepared for the avalanche of data that is going to result from cloud and big data analytics in the coming year and beyond. This data growth will challenge the capabilities of even the most robust backup environments. Short-term solutions, solutions with limited scalability, and solutions that add complexity to the data center will no longer be viable options.

8.2 Continuous re-evaluation of backup strategy

Having a clear backup strategy and re-evaluating it according to the business requirements and needs at least yearly, is necessary to prevent disaster. The following points need to be addressed and re-evaluated every time there is a new Business Impact Analysis on backup and restore.

- Address specific business needs. Define a solution for data backup as part of a consultative process, while addressing all specific business needs.
- Conduct TCO analysis. Before moving to a new data backup service, a TCO (total cost of ownership) analysis should be done to determine the payback period for moving to a new backup service. Use a provider that can integrate archives, so you can move data sets from a backup plan to an archive plan and provides online search and retrieval functionality. And make sure your provider has good management processes, good quality reporting and a secure facility and connectivity.
- Test before provisioning. Put the backup provision in place long before you need it.
- Encrypt backup data. Security is of prime concern in all forms of computing, especially on cloud-based solutions and services. To ensure security and privacy always use encrypted backups.
- Follow governance and compliance rules. Always make sure to follow governance and compliance rules. For example, regulatory compliance related to where data may move or be stored when different countries or regions are involved, or compliance related to retention periods of data. If you are purchasing backup service from vendor and a cloud provider then ensure the backup service vendor is also following the governing laws for that region.
- Bulk data import process governance. Data center staff should be familiar with process and procedures related to bulk data import wherein data is shipped on removable media storage to on-premise. This option will be critical when faster data recovery is needed for large data backups. In addition customer should have in place good governance for data-import process – such as who is authorized to receive the removable storage media and who is notified.
- Backup locally and remotely. Data that is recovered frequently should be backed up to both on-premise storage and to cloud storage. This is because the on-premise backup assures faster recovery, while the off-premise copy can be used for disaster recovery purposes.

References

- Backup locally to ensure public accessibility. If the purpose of putting data in the cloud is for public accessibility, then backup the data locally even before storing in cloud.
- Engage multiple vendors. If one can afford it, it is recommended to backup very important and critical data to multiple vendors to mitigate risks. This gives an extra level of safety in case one of the vendors has an outage at a business-crucial time.
- Ensure data interoperability. Ensure that the backed up data can be recovered on premise and/or to another cloud vendor.

8.3 Stress of Uncertainty

In the end, the stress of uncertainty the last but not least important factor for the future developments of backup. With key information about backups divided on multiple, individual systems, IT staff have more uncertainty to deal with and a harder time making holistic, informed decisions about their company-wide backup environments.

A solution should be chosen that enables IT data managers to get fast, accurate information on the status of their entire backup environment. Robust dashboard functionality can not only enable a single administrator to manage more backup data, it can also enable them to reduce inefficiencies in their backup environment, plan for future capacity needs, and ensure restore Service Level Agreements (SLAs) are achievable.

While the impact on morale is harder to quantify, the human costs of sprawl can quickly affect the bottom line with increased staff turnover, low staff productivity, increased overtime. By consolidating backups and automating backup tasks, companies can free key IT staff for more productive and more gratifying work but most importantly achieve their mission; meeting their SLAs and reducing the business risk of data loss and downtime.

Scalable, enterprise-class systems also provide a tighter level of management control and reporting to ensure IT departments get the most value from their backup investment and more accurate planning for future needs.

References

- [1] Omar H Alhazmi and Yashwant K Malaiya. Evaluating disaster recovery plans using the cloud. In *Reliability and Maintainability Symposium (RAMS), 2013 Proceedings—Annual*, pages 1–6. IEEE, 2013.
- [2] Charlotte Brooks, Matthew Bedernjak, Igor Juran, and John Merryman. *Disaster recovery strategies with Tivoli Storage Management*. IBM Corp., 2002.
- [3] Peter Fallara. Disaster recovery planning. *Potentials, IEEE*, 22(5):42–44, 2003.
- [4] Shon Harris. *CISSP All-in-One Exam Guide*. McGraw-Hill Osborne Media, 6th edition, 2012.

References

- [5] Dave Russel and Pushan Rinnen. Magic quadrant for enterprise backup software and integrated appliances. *Gartner*, June 2014.
- [6] Montri Wiboonratr and Kitti Kosavisutte. Optimal strategic decision for disaster recovery. *International Journal of Management Science and Engineering Management*, 4(4):260–269, 2009.
- [7] Wikipedia. Data deduplication — Wikipedia, the free encyclopedia. http://en.wikipedia.org/w/index.php?title=Data_deduplication&oldid=635040860, 2014. [Online; accessed 23-December-2014].
- [8] Wikipedia. Data integrity — Wikipedia, the free encyclopedia. http://en.wikipedia.org/w/index.php?title=Data_integrity&oldid=636149302, 2014. [Online; accessed 23-December-2014].